PROGRAM GUIDE FRIDAY, JUNE 13

Friday, June 13

7:00 - 18:00 Registration / Badge Pickup (Summit Lobby)
7:00 - 18:00 Registration / Badge Pickup (ExHall Concourse)
7:00 - 17:00 Press Room (203 B)
7:00 - 17:00 Mother's Room (Level 1 near Room 101 and on Level 3 near Exhibit Hall D)
7:00 - 17:00 Prayer or Quiet Room (203 A)

7:00 - 9:00 Breakfast (ExHall C)
8:00 - 8:30 Poster Setup (ExHall D)

8:30 - 9:00 Welcome & Awards (Karl Dean Ballroom)

9:00 - 10:15 Oral Session 1A: Image and Video Synthesis (Karl Dean Ballroom)

2 - Award candidate paper

- Motion Prompting: Controlling Video Generation with Motion Trajectories, Daniel Geng, Charles Herrmann, Junhwa Hur, Forrester Cole, Serena Zhang, Tobias Pfaff, Tatiana Lopez-Guevara, Yusuf Aytar, Michael Rubinstein, Chen Sun, Oliver Wang, Andrew Owens, Deging Sun
- 2 Go-with-the-Flow: Motion-Controllable Video Diffusion Models Using Real-Time Warped Noise, Ryan Burgert, Yuancheng Xu, Wenqi Xian, Oliver Pilarski, Pascal Clausen, Mingming He, Li Ma, Yitong Deng, Lingxiao Li, Mohsen Mousavi, Michael Ryoo, Paul Debevec, Ning Yu
- B LookingGlass: Generative Anamorphoses via Laplacian Pyramid Warping, Pascal Chang, Sergio Sancho, Jingwei Tang, Markus Gross, Vinicius Azevedo
- 4 Alias-Free Latent Diffusion Models: Improving Fractional Shift Equivariance of Diffusion Latent Space, Yifan Zhou, Zeqi Xiao, Shuai Yang, Xingang Pan
- 5 RandAR: Decoder-only Autoregressive Visual Generation in Random Orders, Ziqi Pang, Tianyuan Zhang, Fujun Luan, Yunze Man, Hao Tan, Kai Zhang, William T. Freeman, Yu-Xiong Wang

9:00 - 10:15 Oral Session 1B: Interpretability and Evaluation (ExHall A2)

- OpenING: A Comprehensive Benchmark for Judging Open-ended Interleaved Image-Text Generation, Pengfei Zhou, Xiaopeng Peng, Jiajun Song, Chuanhao Li, Zhaopan Xu, Yue Yang, Ziyao Guo, Hao Zhang, Yuqi Lin, Yefei He, Lirui Zhao, Shuo Liu, Tianhua Li, Yuxuan Xie, Xiaojun Chang, Yu Qiao, Wenqi Shao, Kaipeng Zhang
- 2 LibraGrad: Balancing Gradient Flow for Universally Better Vision Transformer Attributions, Faridoun Mehri, Mahdieh Soleymani Baghshah, Mohammad Taher Pilehvar
- 3 Do We Always Need the Simplicity Bias? Looking for Optimal Inductive Biases in the Wild, Damien Teney, Liangze Jiang, Florin Gogianu, Ehsan Abbasnejad
- Molmo and PixMo: Open Weights and Open Data for State-of-the-Art Vision-Language Models, Matt Deitke, Christopher Clark, Sangho Lee, Rohun Tripathi, Yue Yang, Jae Sung Park, Mohammadreza Salehi, Niklas Muennighoff, Kyle Lo, Luca Soldaini, Jiasen Lu, Taira Anderson, Erin Bransom, Kiana Ehsani, Huong Ngo, YenSung Chen, Ajay Patel, Mark Yatskar, Chris Callison-Burch, Andrew Head, Rose Hendrix, Favyen Bastani, Eli VanderBilt, Nathan Lambert, Yvonne Chou, Arnavi Chheda, Jenna Sparks, Sam Skjonsberg, Michael Schmitz, Aaron Sarnat, Byron Bischoff, Pete Walsh, Chris Newell, Piper Wolters, Tanmay Gupta, Kuo-Hao Zeng, Jon Borchardt, Dirk Groeneveld, Crystal Nam, Sophie Lebrecht, Caitlin Wittlif, Carissa Schoenick, Oscar Michel, Ranjay Krishna, Luca Weihs, Noah A. Smith, Hannaneh Hajishirzi, Ross Girshick, Ali Farhadi, Aniruddha Kembhavi
- 5 Rethinking Vision-Language Model in Face Forensics: Multi-Modal Interpretable Forged Face Detector, Xiao Guo, Xiufeng Song, Yue Zhang, Xiaohong Liu, Xiaoming Liu

9:00 - 10:15 Oral Session 1C: Image Processing and Deep Architectures (Davidson Ballroom)

1 CleanDIFT: Diffusion Features without Noise, Nick Stracke, Stefan Andreas Baumann, Kolja Bauer, Frank Fundel, Björn Ommer

- 2 OverLoCK: An Overview-first-Look-Closely-next ConvNet with Context-Mixing Dynamic Kernels, Meng Lou, Yizhou Yu
- Towards Explicit Geometry-Reflectance Collaboration for Generalized LiDAR Segmentation in Adverse Weather, Longyu Yang, Ping Hu, Shangbo Yuan, Lu Zhang, Jun Liu, Hengtao Shen, Xiaofeng Zhu
- 4 DiffFNO: Diffusion Fourier Neural Operator, Xiaoyi Liu, Hao Tang
- 5 Removing Reflections from RAW Photos, Eric Kee, Adam Pikielny, Kevin Blackburn-Matzen, Marc Levoy

10:00 - 10:30 Poster Setup (ExHall D **10:00 - 11:00 Coffee Break** (ExHall D)

10:30 - 12:30 Poster Session 1 and Exhibit Hall (ExHall D)

* - Highlight paper
□ - Award candidate paper
□ - Oral Paper
☆ - Outstanding Reviewer

- Prosody-Enhanced Acoustic Pre-training and Acoustic-Disentangled Prosody Adapting for Movie Dubbing, Zhedong Zhang, Liang Li, Chenggang Yan, Chunshan Liu, Anton van den Hengel, Yuankai Qi
- Wav2Sem: Plug-and-Play Audio Semantic Decoupling for 3D Speech-Driven Facial Animation, Hao Li, Ju Dai, Xin Zhao, Feng Zhou, Junjun Pan, Lei Li
- 3 Sonic: Shifting Focus to Global Audio Perception in Portrait
- Animation, Xiaozhong Ji, Xiaobin Hu, Zhihong Xu, Junwei Zhu, Chuming Lin, Qingdong He, Jiangning Zhang, Donghao Luo, Yi Chen, Qin Lin, Qinglin Lu, Chengjie Wang
- 4 Towards High-fidelity 3D Talking Avatar with Personalized Dynamic Texture, Xuanchen Li, Jianyu Wang, Yuhao Cheng, Yikun Zeng, Xingyu Ren, Wenhan Zhu, Weiming Zhao, Yichao Yan
- 5 Electromyography-Informed Facial Expression Reconstruction for
- * Physiological-Based Synthesis and Analysis, Tim Büchner,
- Christoph Anders, Orlando Guntinas-Lichius, Joachim Denzler
- 6 High-Fidelity Relightable Monocular Portrait Animation with Lighting-Controllable Video Diffusion Model, *Mingtao Guo, Guanyu Xing, Yanli Liu*
- 7 Quaffure: Real-Time Quasi-Static Neural Hair Simulation, Tuur Stuyck, Gene Wei-Chin Lin, Egor Larionov, Hsiao-yu Chen, Aljaz Bozic, Nikolaos Sarafianos, Doug Roble
- 8 GPAvatar: High-fidelity Head Avatars by Learning Efficient Gaussian Projections, Wei-Qi Feng, Dong Han, Ze-Kang Zhou, Shunkai Li, Xiaoqiang Liu, Pengfei Wan, Di Zhang, Miao Wang
- 9 HERA: Hybrid Explicit Representation for Ultra-Realistic Head Avatars, *Hongrui Cai, Yuting Xiao, Xuan Wang, Jiafei Li, Yudong Guo, Yanbo Fan, Shenghua Gao, Juyong Zhang*
- 10 GASP: Gaussian Avatars with Synthetic Priors, Jack Saunders, Charlie Hewitt, Yanan Jian, Marek Kowalski, Tadas Baltrusaitis, Yiye Chen, Darren Cosker, Virginia Estellers, Nicholas Gydé, Vinay P. Namboodiri, Benjamin E. Lundell
- 11 FRESA: Feedforward Reconstruction of Personalized Skinned
- * Avatars from Few Images, Rong Wang, Fabian Prada, Ziyan Wang, Zhongshi Jiang, Chengxiang Yin, Junxuan Li, Shunsuke Saito, Igor Santesteban, Javier Romero, Rohan Joshi, Hongdong Li, Jason Saragih, Yaser Sheikh
- 2 DAGSM: Disentangled Avatar Generation with GS-enhanced Mesh, Jingyu Zhuang, Di Kang, Linchao Bao, Liang Lin, Guanbin Li
- HumanRig: Learning Automatic Rigging for Humanoid Character in a Large Scale Dataset, Zedong Chu, Feng Xiong, Meiduo Liu, Jinzhi Zhang, Mingqi Shao, Zhaoxu Sun, Di Wang, Mu Xu
- 14 SKDream: Controllable Multi-view and 3D Generation with
- Arbitrary Skeletons, Yuanyou Xu, Zongxin Yang, Yi Yang
 FreeUV: Ground-Truth-Free Realistic Facial UV Texture Recovery
- via Cross-Assembly Inference Strategy, Xingchao Yang, Takafumi Taketomi, Yuki Endo, Yoshihiro Kanamori
- MultiGO: Towards Multi-level Geometry Learning for Monocular 3D Textured Human Reconstruction, Gangjian Zhang, Nanjie Yao, Shunsi Zhang, Hanfeng Zhao, Guoliang Pang, Jian Shu, Hao Wang
- 17 GaussianIP: Identity-Preserving Realistic 3D Human Generation via Human-Centric Diffusion Prior, Zichen Tang, Yuan Yao, Miaomiao Cui, Liefeng Bo, Hongyu Yang

PROGRAM GUIDE FRIDAY, JUNE 13

- Shining Yourself: High-Fidelity Ornaments Virtual Try-on with Diffusion Model, Yingmao Miao, Zhanpeng Huang, Rui Han, Zibin Wang, Chenhao Lin, Chao Shen
- SynthLight: Portrait Relighting with Diffusion Model by Learning to Re-render Synthetic Faces, Sumit Chaturvedi, Mengwei Ren, Yannick Hold-Geoffroy, Jingyuan Liu, Julie Dorsey, Zhixin Shu
- Comprehensive Relighting: Generalizable and Consistent Monocular Human Relighting and Harmonization, Junying Wang, Jingyuan Liu, Xin Sun, Krishna Kumar Singh, Zhixin Shu, He Zhang, Jimei Yang, Nanxuan Zhao, Tuanfeng Y. Wang, Simon S. Chen, Ulrich Neumann, Jae Shin Yoon
- Polarized Color Screen Matting, Kenji Enomoto, Scott Cohen,
- Brian Price, TJ Rhodes
- SLVR: Super-Light Visual Reconstruction via Blueprint Controllable Convolutions and Exploring Feature Diversity Representation, Ning Ni, Libao Zhang
- Proximal Algorithm Unrolling: Flexible and Efficient Reconstruction Networks for Single-Pixel Imaging, Ping Wang, Lishun Wang, Gang Qu, Xiaodong Wang, Yulun Zhang, Xin Yuan
- Glossy Object Reconstruction with Cost-effective Polarized
- Acquisition, Bojian Wu, Yifan Peng, Ruizhen Hu, Xiaowei Zhou
- Blurry-Edges: Photon-Limited Depth Estimation from Defocused Boundaries, Wei Xu, Charles James Wagner, Junjie Luo, Qi Guo
- LumiNet: Latent Intrinsics Meets Diffusion Models for Indoor Scene Relighting, Xiaoyan Xing, Konrad Groh, Sezer Karaoglu, Theo Gevers, Anand Bhattad
- LEDiff: Latent Exposure Diffusion for HDR Generation, Chao Wang, Zhihao Xia, Thomas Leimkuhler, Karol Myszkowski, Xuaner Zhang
- IRIS: Inverse Rendering of Indoor Scenes from Low Dynamic Range Images, Chih-Hao Lin, Jia-Bin Huang, Zhengqin Li, Zhao Dong, Christian Richardt, Tuotuo Li, Michael Zollhöfer, Johannes Kopf, Shenlong Wang, Changil Kim
- Differentiable Inverse Rendering with Interpretable Basis BRDFs, Hoon-Gyu Chung, Seokjun Choi, Seung-Hwan Baek
- Hardware-Rasterized Ray-Based Gaussian Splatting, Samuel Rota Bulò,
- Nemanja Bartolovic, Lorenzo Porzi, Peter Kontschieder
- TensoFlow: Tensorial Flow-based Sampler for Inverse Rendering, Chun Gu, Xiaofei Wei, Li Zhang, Xiatian Zhu
- LIRM: Large Inverse Rendering Model for Progressive Reconstruction of Shape, Materials and View-dependent Radiance Fields, Zhengqin Li, Dilin Wang, Ka Chen, Zhaoyang Lv, Thu Nguyen-Phuoc, Milim Lee, Jia-Bin Huang, Lei Xiao, Yufeng Zhu, Carl S. Marshall, Yuheng Ren, Richard Newcombe, Zhao Dong
- Gaussian Splashing: Unified Particles for Versatile Motion Synthesis and Rendering, Yutao Feng, Xiang Feng, Yintong Shang, Ying Jiang, Chang Yu, Zeshun Zong, Tianjia Shao, Hongzhi Wu, Kun Zhou, Chenfanfu Jiang, Yin Yang
- Accurate Differential Operators for Hybrid Neural Fields, Aditya Chetan, Guandao Yang, Zichen Wang, Steve Marschner, Bharath Hariharan
- Learning Extremely High Density Crowds as Active Matters, Feixiang He, Jiangbei Yue, Jialin Zhu, Armin Seyfried, Dan Casas, Julien Pettré, He Wang
- TexGaussian: Generating High-quality PBR Material via Octree-based 3D Gaussian Splatting, Bojun Xiong, Jialun Liu, Jiakui Hu, Chenming Wu, Jinbo Wu, Xing Liu, Chen Zhao, Errui Ding, Zhouhui Lian
- Real-time Free-view Human Rendering from Sparse-view RGB Videos using Double Unprojected Textures, Guoxing Sun, Rishabh Dabral, Heming Zhu, Pascal Fua, Christian Theobalt,
- Marc Habermann RoomPainter: View-Integrated Diffusion for Consistent Indoor Scene Texturing, Zhipeng Huang, Wangbo Yu, Xinhua Cheng,
- Chengshu Zhao, Yunyang Ge, Mingyi Guo, Li Yuan, Yonghong Tian MVPaint: Synchronized Multi-View Diffusion for Painting Anything 3D, Wei Cheng, Juncheng Mu, Xianfang Zeng, Xin Chen, Angi Pang, Chi Zhang, Zhibin Wang, Bin Fu, Gang Yu, Ziwei Liu, Liang Pan
- Fancy123: One Image to High-Quality 3D Mesh Generation via Plug-and-Play Deformation, Qiao Yu, Xianzhi Li, Yuan Tang, Xu Han, Long Hu, Yixue Hao, Min Chen

ShapeShifter: 3D Variations Using Multiscale and Sparse Point-

- Voxel Diffusion, Nissim Maruani, Wang Yifan, Matthew Fisher, Pierre Alliez, Mathieu Desbrun
- MeshArt: Generating Articulated Meshes with Structure-Guided Transformers, Daoyi Gao, Yawar Siddiqui, Lei Li, Angela Dai
- SceneFactor: Factored Latent 3D Diffusion for Controllable 3D Scene Generation, Aleksey Bokhovkin, Quan Meng, Shubham Tulsiani, Angela Dai
- PrEditor3D: Fast and Precise 3D Shape Editing, Ziya Erkoç, Can Gümeli, Chaoyang Wang, Matthias Nießner, Angela Dai, Peter Wonka, Hsin-Ying Lee, Peiye Zhuang
- LT3SD: Latent Trees for 3D Scene Diffusion, Quan Meng, Lei Li, Matthias Nießner, Angela Dai
- iSegMan: Interactive Segment-and-Manipulate 3D Gaussians, Yian Zhao, Wanshi Xu, Ruochong Zheng, Pengchong Qiao, Chang Liu, Jie Chen
- LOD-GS: Achieving Levels of Detail using Scalable Gaussian Soup, Jianxiong Shen, Yue Qian, Xiaohang Zhan
- MaskGaussian: Adaptive 3D Gaussian Representation from Probabilistic Masks, Yifei Liu, Zhihang Zhong, Yifan Zhan, Sheng Xu, Xiao Sun
- NTR-Gaussian: Nighttime Dynamic Thermal Reconstruction with 4D Gaussian Splatting Based on Thermodynamics, Kun Yang, Yuxiang Liu, Zeyu Cui, Yu Liu, Maojun Zhang, Shen Yan, Qing Wang
- DropoutGS: Dropping Out Gaussians for Better Sparse-view Rendering, Yexing Xu, Longguang Wang, Minglin Chen, Sheng Ao, Li Li, Yulan Guo
- S2Gaussian: Sparse-View Super-Resolution 3D Gaussian Splatting, Yecong Wan, Mingwen Shao, Yuanshuo Cheng, Wangmeng Zuo
- 52 DeSplat: Decomposed Gaussian Splatting for Distractor-Free
- Rendering, Yihao Wang, Marcus Klasson, Matias Turkulainen, Shuzhe Wang, Juho Kannala, Arno Solin
- 53 Neural Hierarchical Decomposition for Single Image Plant Modeling, Zhihao Liu, Zhanglin Cheng, Naoto Yokoya
- 54 Symmetry Strikes Back: From Single-Image Symmetry Detection
- to 3D Generation, Xiang Li, Zixuan Huang, Anh Thai, James M. Reha
- Digital Twin Catalog: A Large-Scale Photorealistic 3D Object
- Digital Twin Dataset, Zhao Dong, Ka Chen, Zhaoyang Lv, Hong-Xing Yu, Yunzhi Zhang, Cheng Zhang, Yufeng Zhu, Stephen Tian, Zhengqin Li, Geordie Moffatt, Sean Christofferson, James Fort, Xiaging Pan, Mingfei Yan, Jiajun Wu, Carl Yuheng Ren, Richard Newcombe
- Zero-Shot Novel View and Depth Synthesis with Multi-View Geometric Diffusion, Vitor Guizilini, Muhammad Zubair Irshad, Dian Chen, Greg Shakhnarovich, Rares Ambrus
- NVComposer: Boosting Generative Novel View Synthesis with Multiple Sparse and Unposed Images, Lingen Li, Zhaoyang Zhang, Yaowei Li, Jiale Xu, Wenbo Hu, Xiaoyu Li, Weihao Cheng, Jinwei Gu, Tianfan Xue, Ying Shan
- HybridGS: Decoupling Transients and Statics with 2D and 3D Gaussian Splatting, Jingyu Lin, Jiaqi Gu, Lubin Fan, Bojian Wu, Yujing Lou, Renjie Chen, Ligang Liu, Jieping Ye
- Wonderland: Navigating 3D Scenes from a Single Image, Hanwen Liang, Junli Cao, Vidit Goel, Guocheng Qian, Sergei Korolev, Demetri Terzopoulos, Konstantinos N. Plataniotis, Sergey Tulyakov, Jian Ren
- SpatialDreamer: Self-supervised Stereo Video Synthesis from Monocular Input, Zhen Lv, Yangqi Long, Congzhentao Huang, Cao Li, Chengfei Lv, Hao Ren, Dian Zheng
- StreetCrafter: Street View Synthesis with Controllable Video Diffusion Models, Yunzhi Yan, Zhen Xu, Haotong Lin, Haian Jin, Haoyu Guo, Yida Wang, Kun Zhan, Xianpeng Lang, Hujun Bao, Xiaowei Zhou, Sida Peng
- DroneSplat: 3D Gaussian Splatting for Robust 3D Reconstruction
- from In-the-Wild Drone Imagery, Jiadong Tang, Yu Gao, Dianyi Yang, Liqi Yan, Yufeng Yue, Yi Yang
- IndoorGS: Geometric Cues Guided Gaussian Splatting for Indoor Scene Reconstruction, Cong Ruan, Yuesong Wang, Tao Guan, Bin Zhang, Lili Ju
- MAC-Ego3D: Multi-Agent Gaussian Consensus for Real-Time Collaborative Ego-Motion and Photorealistic 3D

PROGRAM GUIDE FRIDAY, JUNE 13

- Reconstruction, Xiaohao Xu, Feng Xue, Shibo Zhao, Yike Pan, Sebastian Scherer, Xiaonan Huang
- 65 ShowMak3r: Compositional TV Show Reconstruction, Sangmin Kim, Seunguk Do, Jaesik Park
- 66 4DGC: Rate-Aware 4D Gaussian Compression for Efficient
- Streamable Free-Viewpoint Video, Qiang Hu, Zihan Zheng, Houqiang Zhong, Sihua Fu, Li Song, Xiaoyun Zhang, Guangtao Zhai, Yanfeng Wang
- 67 HiMoR: Monocular Deformable Gaussian Reconstruction with Hierarchical Motion Representation, Yiming Liang, Tianhan Xu, Yuta Kikuchi
- 68 EnliveningGS: Active Locomotion of 3DGS, Siyuan Shen, Tianjia Shao, Kun Zhou, Chenfanfu Jiang, Yin Yang
- 69 HOP: Heterogeneous Topology-based Multimodal Entanglement for Co-Speech Gesture Generation, Hongye Cheng, Tianyu Wang, Guangsi Shi, Zexing Zhao, Yanwei Fu
- 70 Stable-SCore: A Stable Registration-based Framework for 3D Shape Correspondence, Haolin Liu, Xiaohang Zhan, Zizheng Yan, Zhongjin Luo, Yuxin Wen, Xiaoguang Han
- 71 Active Hyperspectral Imaging Using an Event Camera,
- * Bohan Yu, Jinxiu Liang, Zhuofeng Wang, Bin Fan, Art Subpa-asa, Boxin Shi, Imari Sato
- 72 SphereUFormer: A U-Shaped Transformer for Spherical 360 Perception, *Yaniv Benny, Lior Wolf*
- 73 Decoupling Fine Detail and Global Geometry for Compressed Depth Map Super-Resolution, Huan Zheng, Wencheng Han, Jianbing Shen
- 74 A Unified Image-Dense Annotation Generation Model for Underwater Scenes, *Hongkai Lin, Dingkang Liang, Zhenghao Qi, Xiang Bai*
- 75 Active Event-based Stereo Vision, *Jianing Li, Yunjian Zhang, Haiqian Han, Xiangyang Ji*
- 76 PanDA: Towards Panoramic Depth Anything with Unlabeled Panoramas and Mobius Spatial Augmentation, Zidong Cao, Jinjing Zhu, Weiming Zhang, Hao Ai, Haotian Bai, Hengshuang Zhao, Lin Wang
- 77 Flow-NeRF: Joint Learning of Geometry, Poses, and Dense Flow within Unified Neural Representations. *Xunzhi Zheng, Dan Xu*
- 78 OmniStereo: Real-time Omnidireactional Depth Estimation with Multiview Fisheye Cameras, *Jiaxi Deng, Yushen Wang, Haitao Meng, Zuoxun Hou, Yi Chang, Gang Chen*
- 79 Stereo Anywhere: Robust Zero-Shot Deep Stereo Matching Even
- ☆ Where Either Stereo or Mono Fail, Luca Bartolomei, Fabio Tosi, Matteo Poggi, Stefano Mattoccia
- 80 UniK3D: Universal Camera Monocular 3D Estimation,
- Luigi Piccinelli, Christos Sakaridis, Mattia Segu, Yung-Hsu Yang, Siyuan Li. Wim Abbeloos. Luc Van Gool
- 81 Structure-from-Motion with a Non-Parametric Camera Model,
- * Yihan Wang, Linfei Pan, Marc Pollefeys, Viktor Larsson
- 82 MUSt3R: Multi-view Network for Stereo 3D
- * Reconstruction, Yohann Cabon, Lucas Stoffl, Leonid Antsfeld, Gabriela Csurka, Boris Chidlovskii, Jerome Revaud, Vincent Leroy
- 83 Extreme Rotation Estimation in the Wild, Hana Bezalel, Dotan Ankri, Ruojin Cai, Hadar Averbach-Elor
- 84 Pow3R: Empowering Unconstrained 3D Reconstruction with Camera and Scene Priors, Wonbong Jang, Philippe Weinzaepfel, Vincent Leroy, Lourdes Agapito, Jerome Revaud
- 85 Gaussian Splatting Feature Fields for (Privacy-Preserving) Visual Localization, Maxime Pietrantoni, Gabriela Csurka, Torsten Sattler
- Dense Match Summarization for Faster Two-view Estimation,

 Jonathan Astermark, Anders Heyden, Viktor Larsson
- 87 Cross-View Completion Models are Zero-shot Correspondence
- * Estimators, Honggyu An, Jin Hyeon Kim, Seonghoon Park, Jaewoo Jung, Jisang Han, Sunghwan Hong, Seungryong Kim
- 88 Uni4D: Unifying Visual Foundation Models for 4D Modeling from
- * a Single Video, David Yifan Yao, Albert J. Zhai, Shenlong Wang
- 89 EquiPose: Exploiting Permutation Equivariance for Relative Camera Pose Estimation, *Yuzhen Liu, Qiulei Dong*
- 90 SemAlign3D: Semantic Correspondence between RGB-Images through Aligning 3D Object-Class Representations, Krispin Wandel, Hesheng Wang

91 PromptHMR: Promptable Human Mesh Recovery, Yufu Wang, Yu Sun, Priyanka Patel, Kostas Daniilidis, Michael J. Black, Muhammed Kocabas

- 92 DynPose: Largely Improving the Efficiency of Human Pose Estimation by a Simple Dynamic Framework, *Yalong Xu, Lin Zhao, Chen Gong, Guangyu Li, Di Wang, Nannan Wang*
- 93 Rethinking Correspondence-based Category-Level Object Pose Estimation, *Huan Ren, Wenfei Yang, Shifeng Zhang, Tianzhu Zhang*
- 94 UA-Pose: Uncertainty-Aware 6D Object Pose Estimation and
- Online Object Completion with Partial References, Ming-Feng Li, Xin Yang, Fu-En Wang, Hritam Basak, Yuyin Sun, Shreekant Gayaka, Min Sun, Cheng-Hao Kuo
- 95 PlanarSplatting: Accurate Planar Surface Reconstruction in
- * 3 Minutes, Bin Tan, Rui Yu, Yujun Shen, Nan Xue
- 96 Prior-free 3D Object Tracking, Xiuqiang Song, Li Jin,
- * Zhengxian Zhang, Jiachen Li, Fan Zhong, Guofeng Zhang, Xueying Qin
- 97 Progressive Correspondence Regenerator for Robust 3D Registration, *Guiyu Zhao, Sheng Ao, Ye Zhang, Kai Xu, Yulan Guo*
- 98 CaMuViD: Calibration-Free Multi-View Detection, Amir Etefaghi Daryani, M. Usman Maqbool Bhutta, Byron Hernandez, Henry Medeiros
- 99 A New Statistical Model of Star Speckles for Learning to Detect and Characterize Exoplanets in Direct Imaging Observations, Théo Bodrito, Olivier Flasseur, Julien Mairal, Jean Ponce, Maud Langlois, Anne-Marie Lagrange
- 100 AG-VPReID: A Challenging Large-Scale Benchmark for Aerial-Ground Video-based Person Re-Identification, Huy Nguyen, Kien Nguyen, Akila Pemasiri, Feng Liu, Sridha Sridharan, Clinton Fookes
- 101 MambaVO: Deep Visual Odometry Based on Sequential Matching Refinement and Training Smoothing, Shuo Wang, Wanting Li, Yongcai Wang, Zhaoxin Fan, Zhe Huang, Xudong Cai, Jian Zhao, Deying Li
- 102 Point-Cache: Test-time Dynamic and Hierarchical Cache for
- Robust and Generalizable Point Cloud Analysis, Hongyu Sun, Qiuhong Ke, Ming Cheng, Yongcai Wang, Deying Li, Chenhui Gou, Jianfei Cai
- 103 HotSpot: Signed Distance Function Optimization with an
- * Asymptotically Sufficient Condition, Zimo Wang, Cheng Wang, Taiki Yoshino, Sirui Tao, Ziyang Fu, Tzu-Mao Li
- 104 High-quality Point Cloud Oriented Normal Estimation via Hybrid Angular and Euclidean Distance Encoding, Yuanqi Li, Jingcheng Huang, Hongshen Wang, Peiyuan Lv, Yansong Liu, Jiuming Zheng, Jie Guo, Yanwen Guo
- 105 A Lightweight UDF Learning Framework for 3D Reconstruction
- Based on Local Shape Functions, Jiangbei Hu, Yanggeng Li, Fei Hou, Junhui Hou, Zhebin Zhang, Shengfa Wang, Na Lei, Ying He
- 106 GenPC: Zero-shot Point Cloud Completion via 3D Generative Priors, An Li, Zhe Zhu, Mingqiang Wei
- 107 UniPre3D: Unified Pre-training of 3D Point Cloud Models with Cross-Modal Gaussian Splatting, Ziyi Wang, Yanran Zhang, Jie Zhou, Jiwen Lu
- 108 DeepLA-Net: Very Deep Local Aggregation Networks for Point Cloud Analysis, Ziyin Zeng, Mingyue Dong, Jian Zhou, Huan Qiu, Zhen Dong, Man Luo, Bijun Li
- SAMBLE: Shape-Specific Point Cloud Sampling for an Optimal Trade-Off Between Local Detail and Global Uniformity, Chengzhi Wu, Yuxin Wan, Hao Fu, Julius Pfrommer, Zeyun Zhong, Junwei Zheng, Jiaming Zhang, Jürgen Beyerer
- 110 PO3AD: Predicting Point Offsets toward Better 3D Point Cloud Anomaly Detection, Jianan Ye, Weiguang Zhao, Xi Yang, Guangliang Cheng, Kaizhu Huang
- HeMoRa: Unsupervised Heuristic Consensus Sampling for Robust Point Cloud Registration, Shaocheng Yan, Yiming Wang, Kaiyan Zhao, Pengcheng Shi, Zhenjun Zhao, Yongjun Zhang, Jiayuan Li
- 112 LogoSP: Local-global Grouping of Superpoints for Unsupervised Semantic Segmentation of 3D Point Clouds, Zihui Zhang, Weisheng Dai, Hongtao Wen, Bo Yang
- 113 AirRoom: Objects Matter in Room Reidentification, Runmao Yao, Yi Du, Zhuoqun Chen, Haoze Zheng, Chen Wang